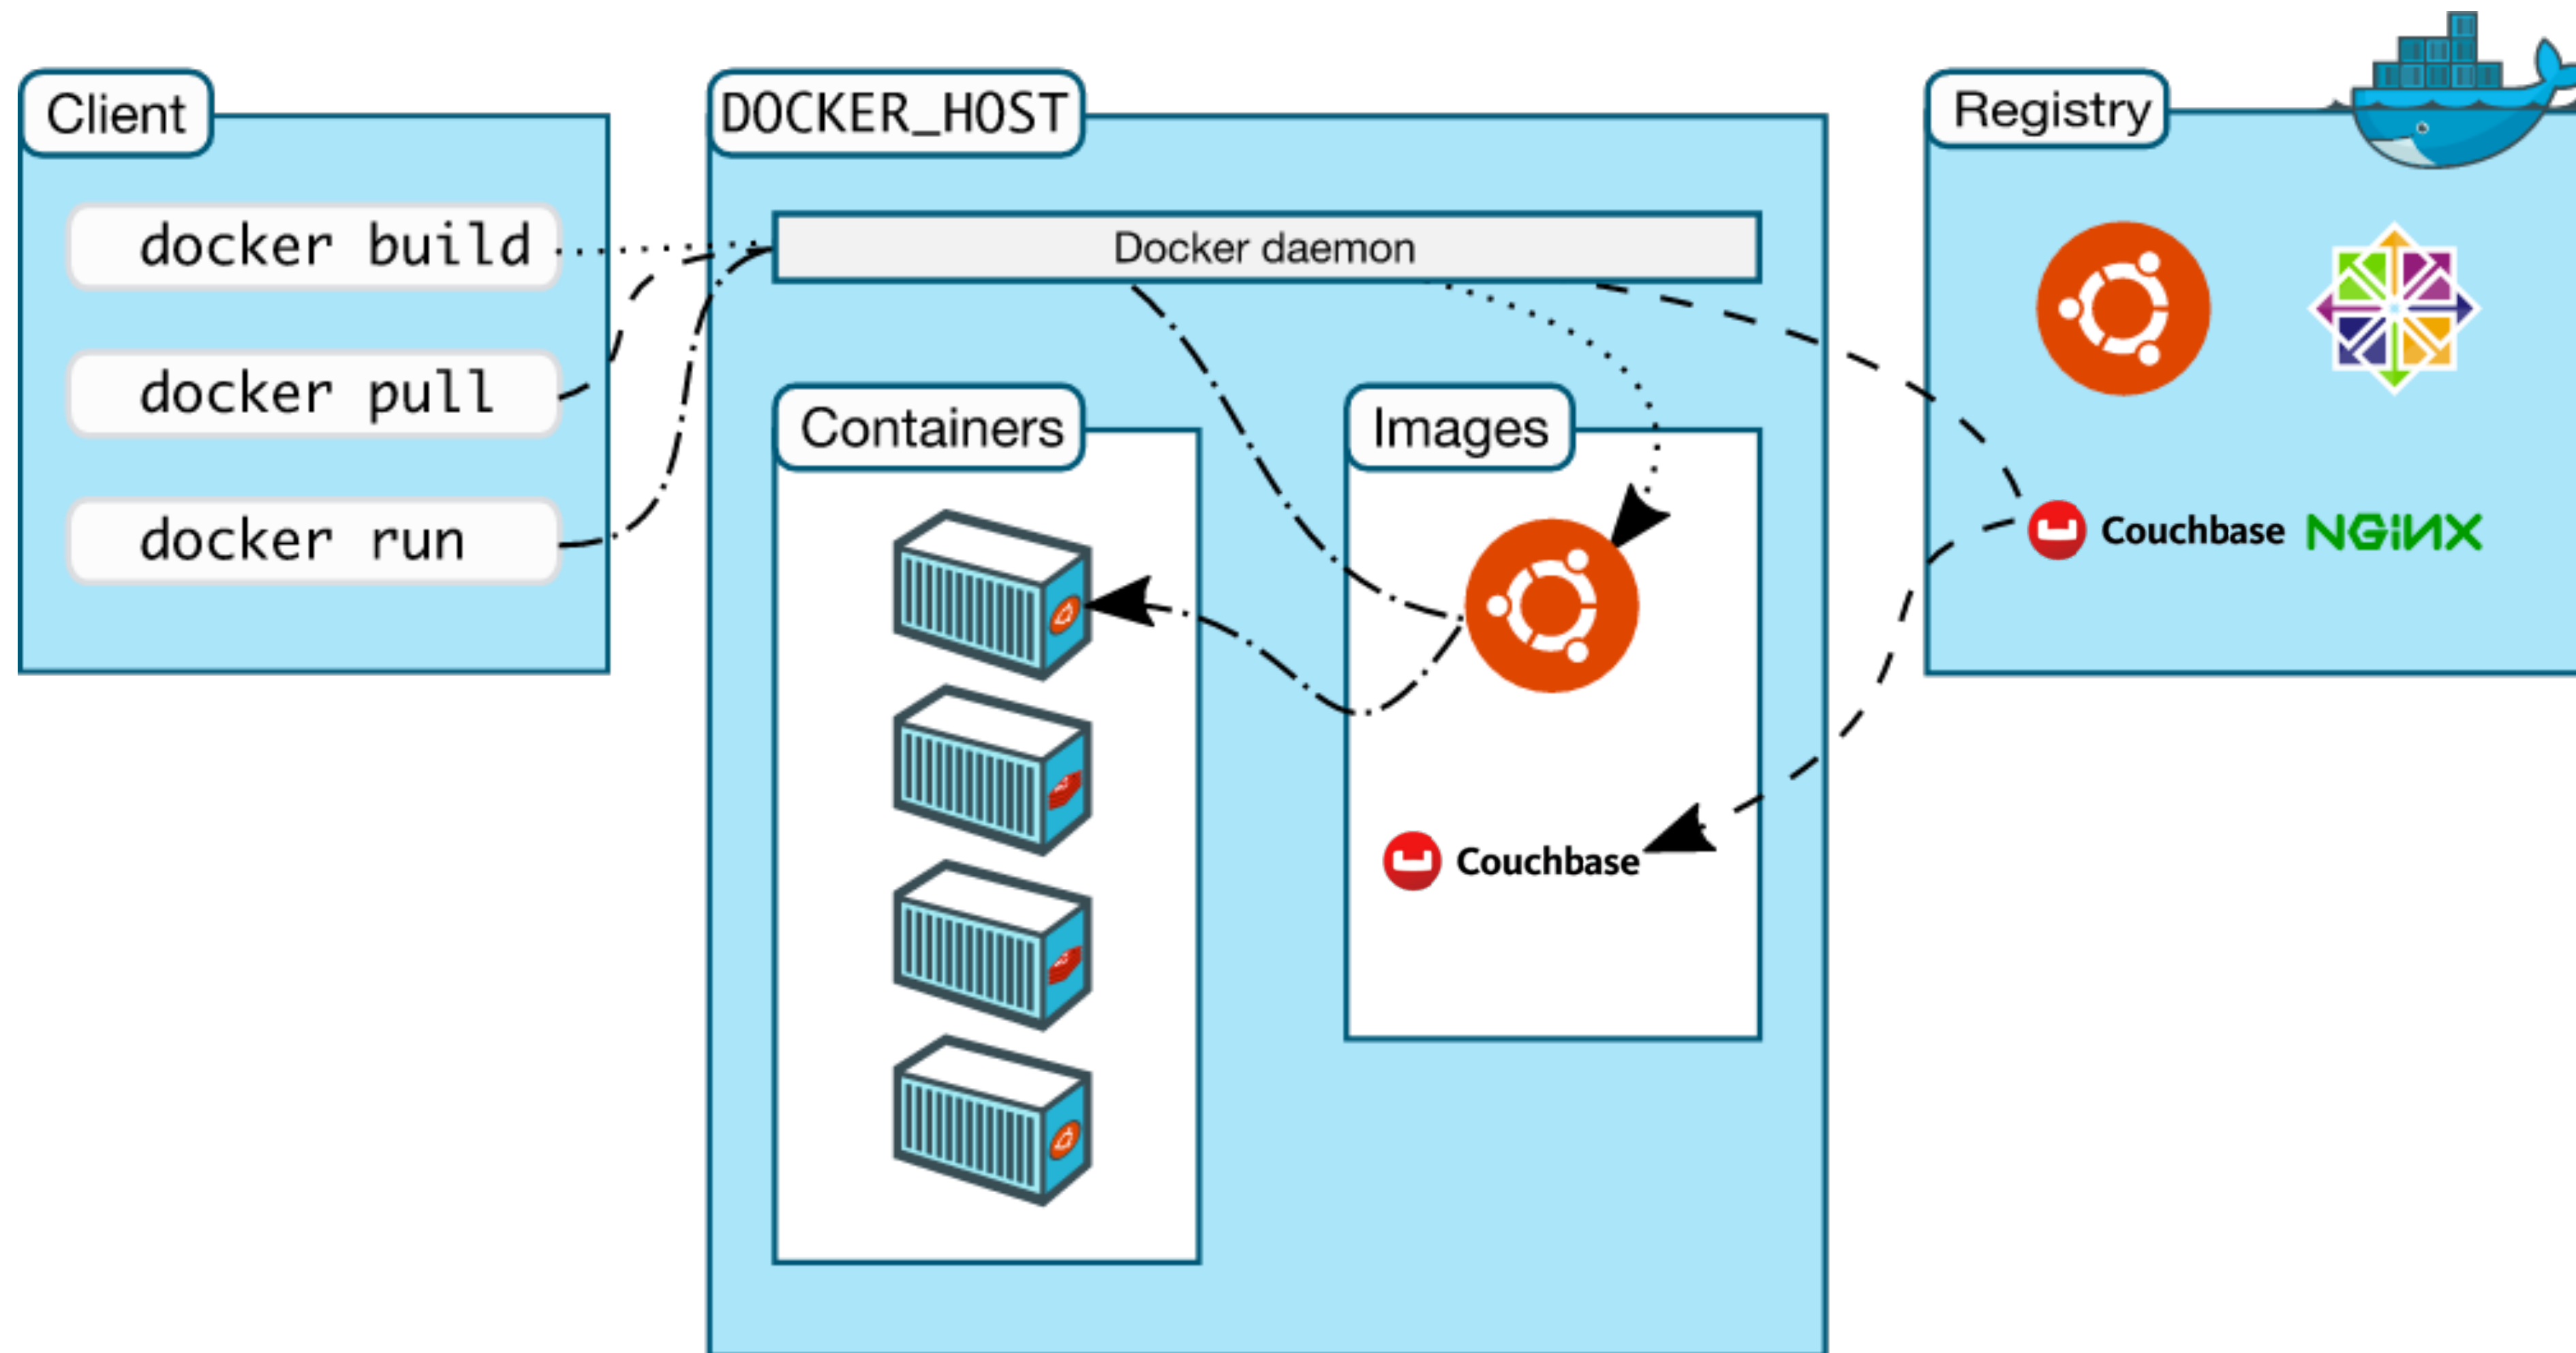



# Container Orchestration on Amazon Web Services

Arun Gupta, @arungupta

# Docker Workflow




# Container Orchestration



# Development using Docker

- Docker Community Edition
  - Docker for Mac/Windows/Linux
  - Monthly edge and quarterly stable releases
  - Native desktop or cloud provider experience




# Swarm-mode: Initialize



```
docker swarm init --listen-addr <ip>:2377 --secret <SECRET>
```

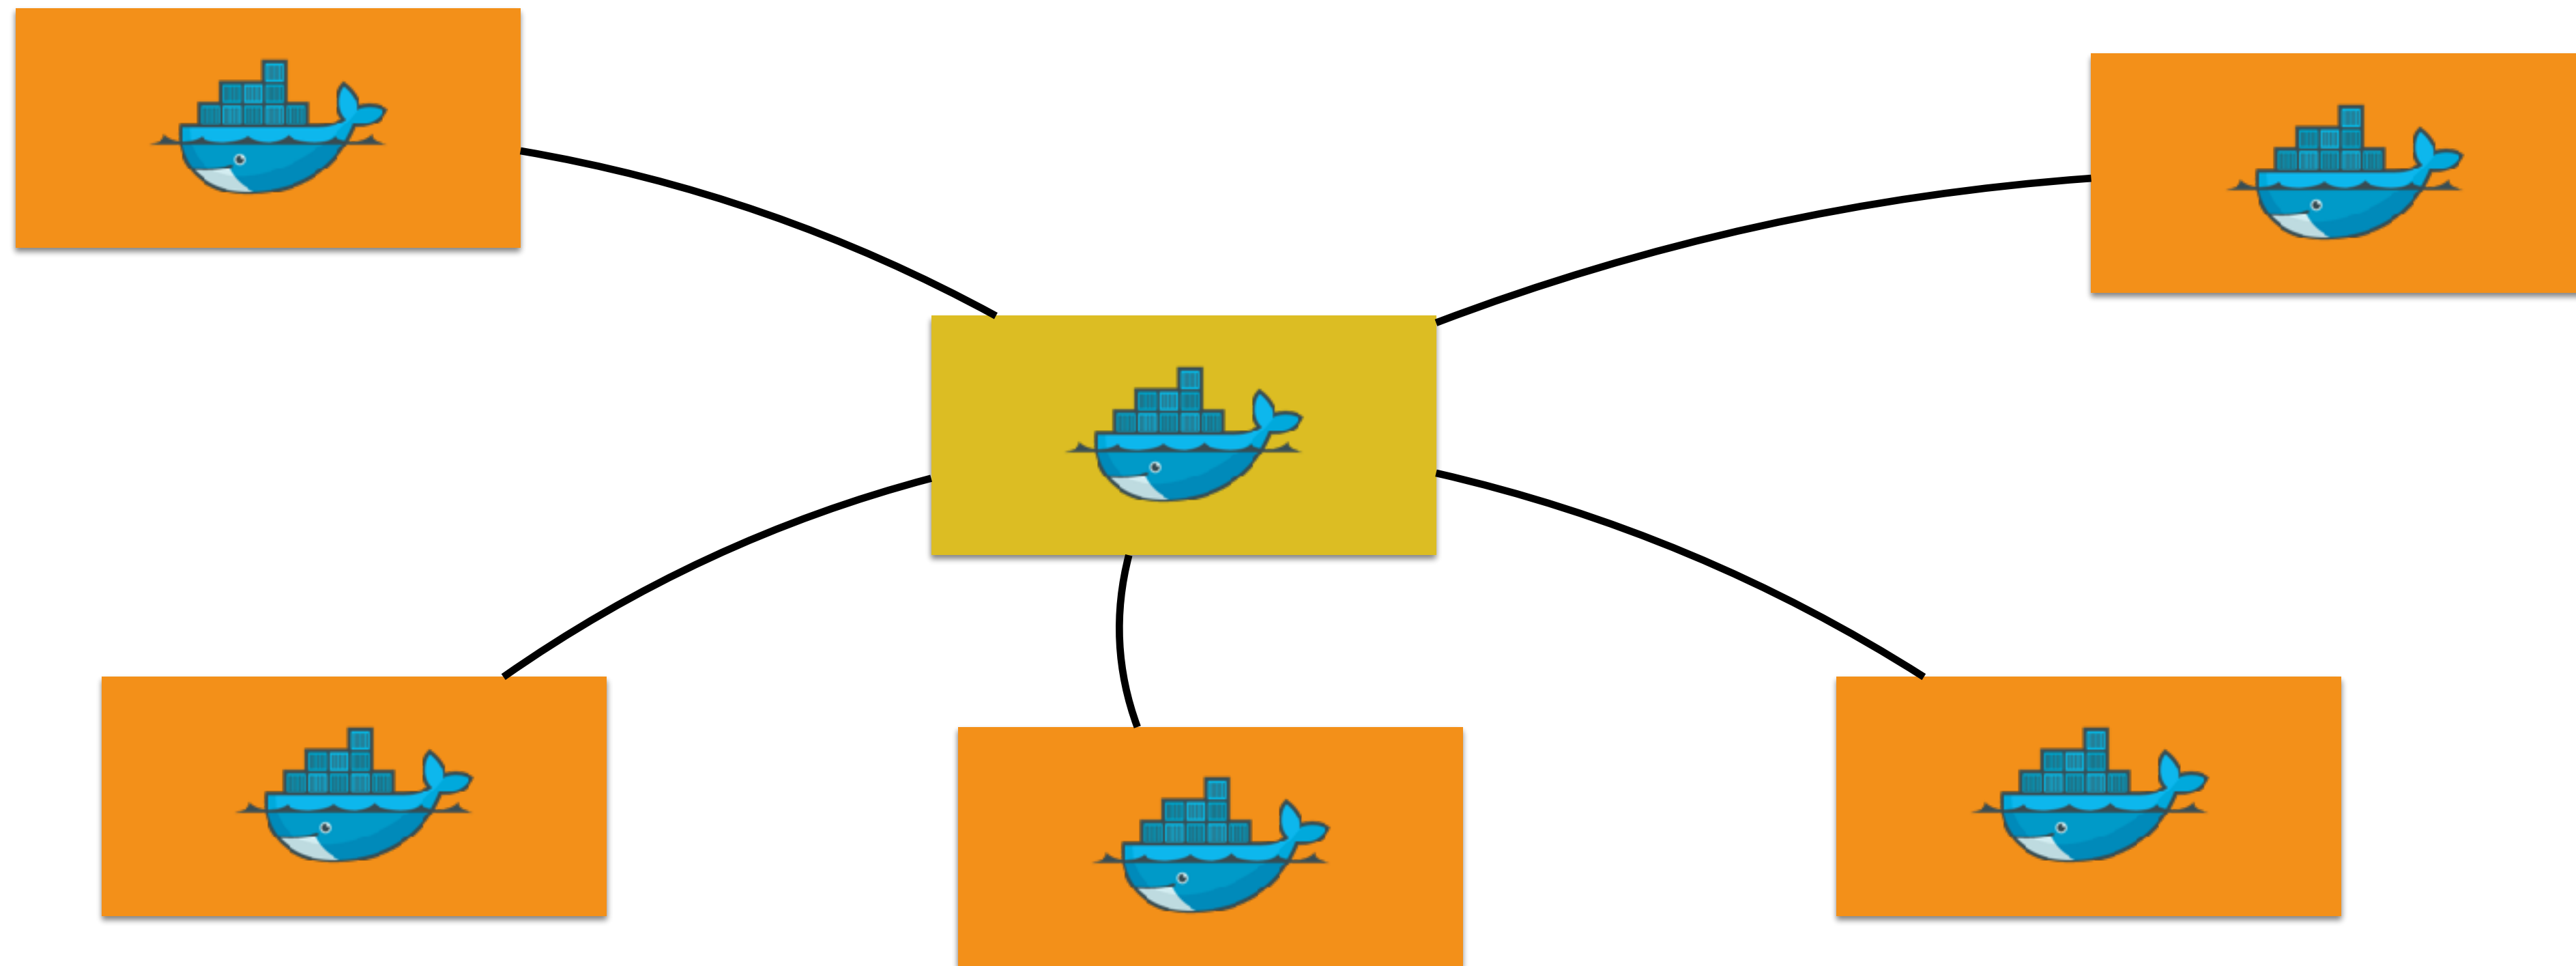
# Swarm-mode: Add Worker



```
docker swarm join --secret <SECRET> <manager>:2377
```




# Swarm-mode: Add More Workers



```
docker swarm join --secret <SECRET> <manager>:2377
```


# Swarm-mode: Primary/Secondary Master



```
docker swarm join --manager --secret <SECRET> --listen-addr  
<master2>:2377 <master1>:2377
```



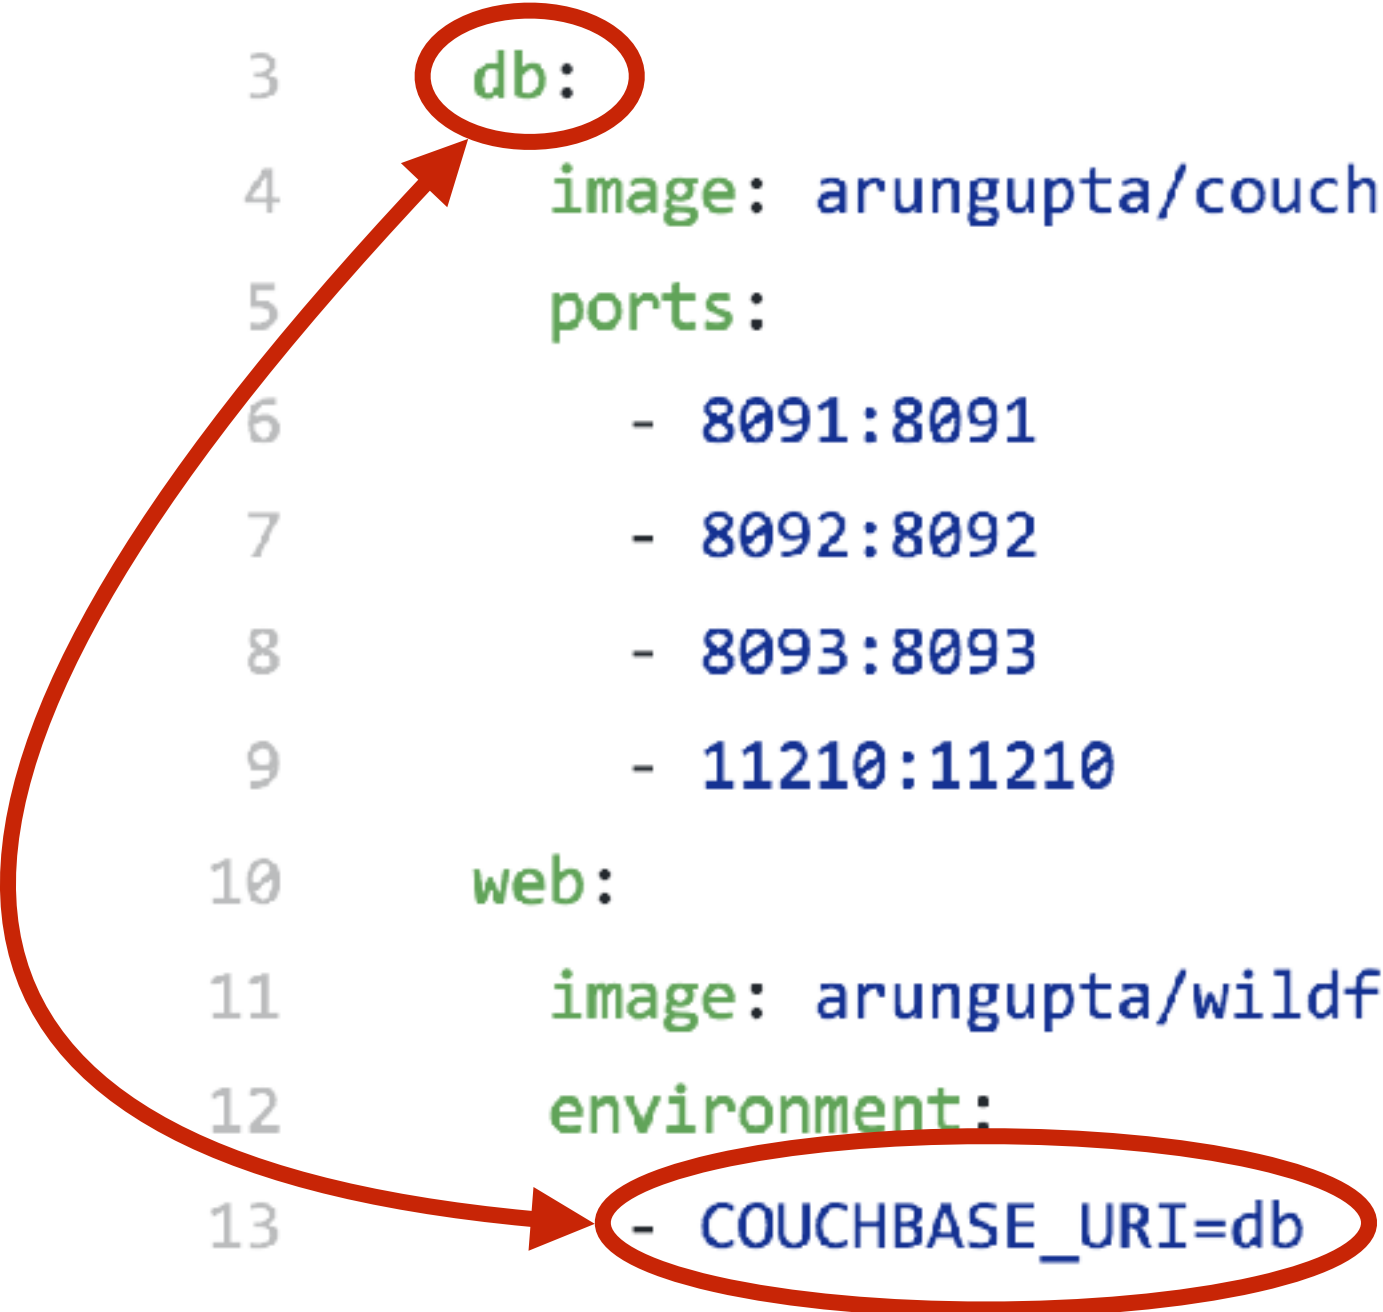
# Swarm-mode: Replicated Service



```
docker service create --replicas 3 --name web jboss/wildfly
```

# Service Discovery with Docker


```
1  version: "3"
2  services:
3    db:
4      image: arungupta/couchbase:travel
5      ports:
6        - 8091:8091
7        - 8092:8092
8        - 8093:8093
9        - 11210:11210
10   web:
11     image: arungupta/wildfly-couchbase-javaee:travel
12     environment:
13       - COUCHBASE_URI=db
14     ports:
15       - 8080:8080
16       - 9990:9990
```



```
docker stack deploy --compose-file=docker-compose.yml webapp
```

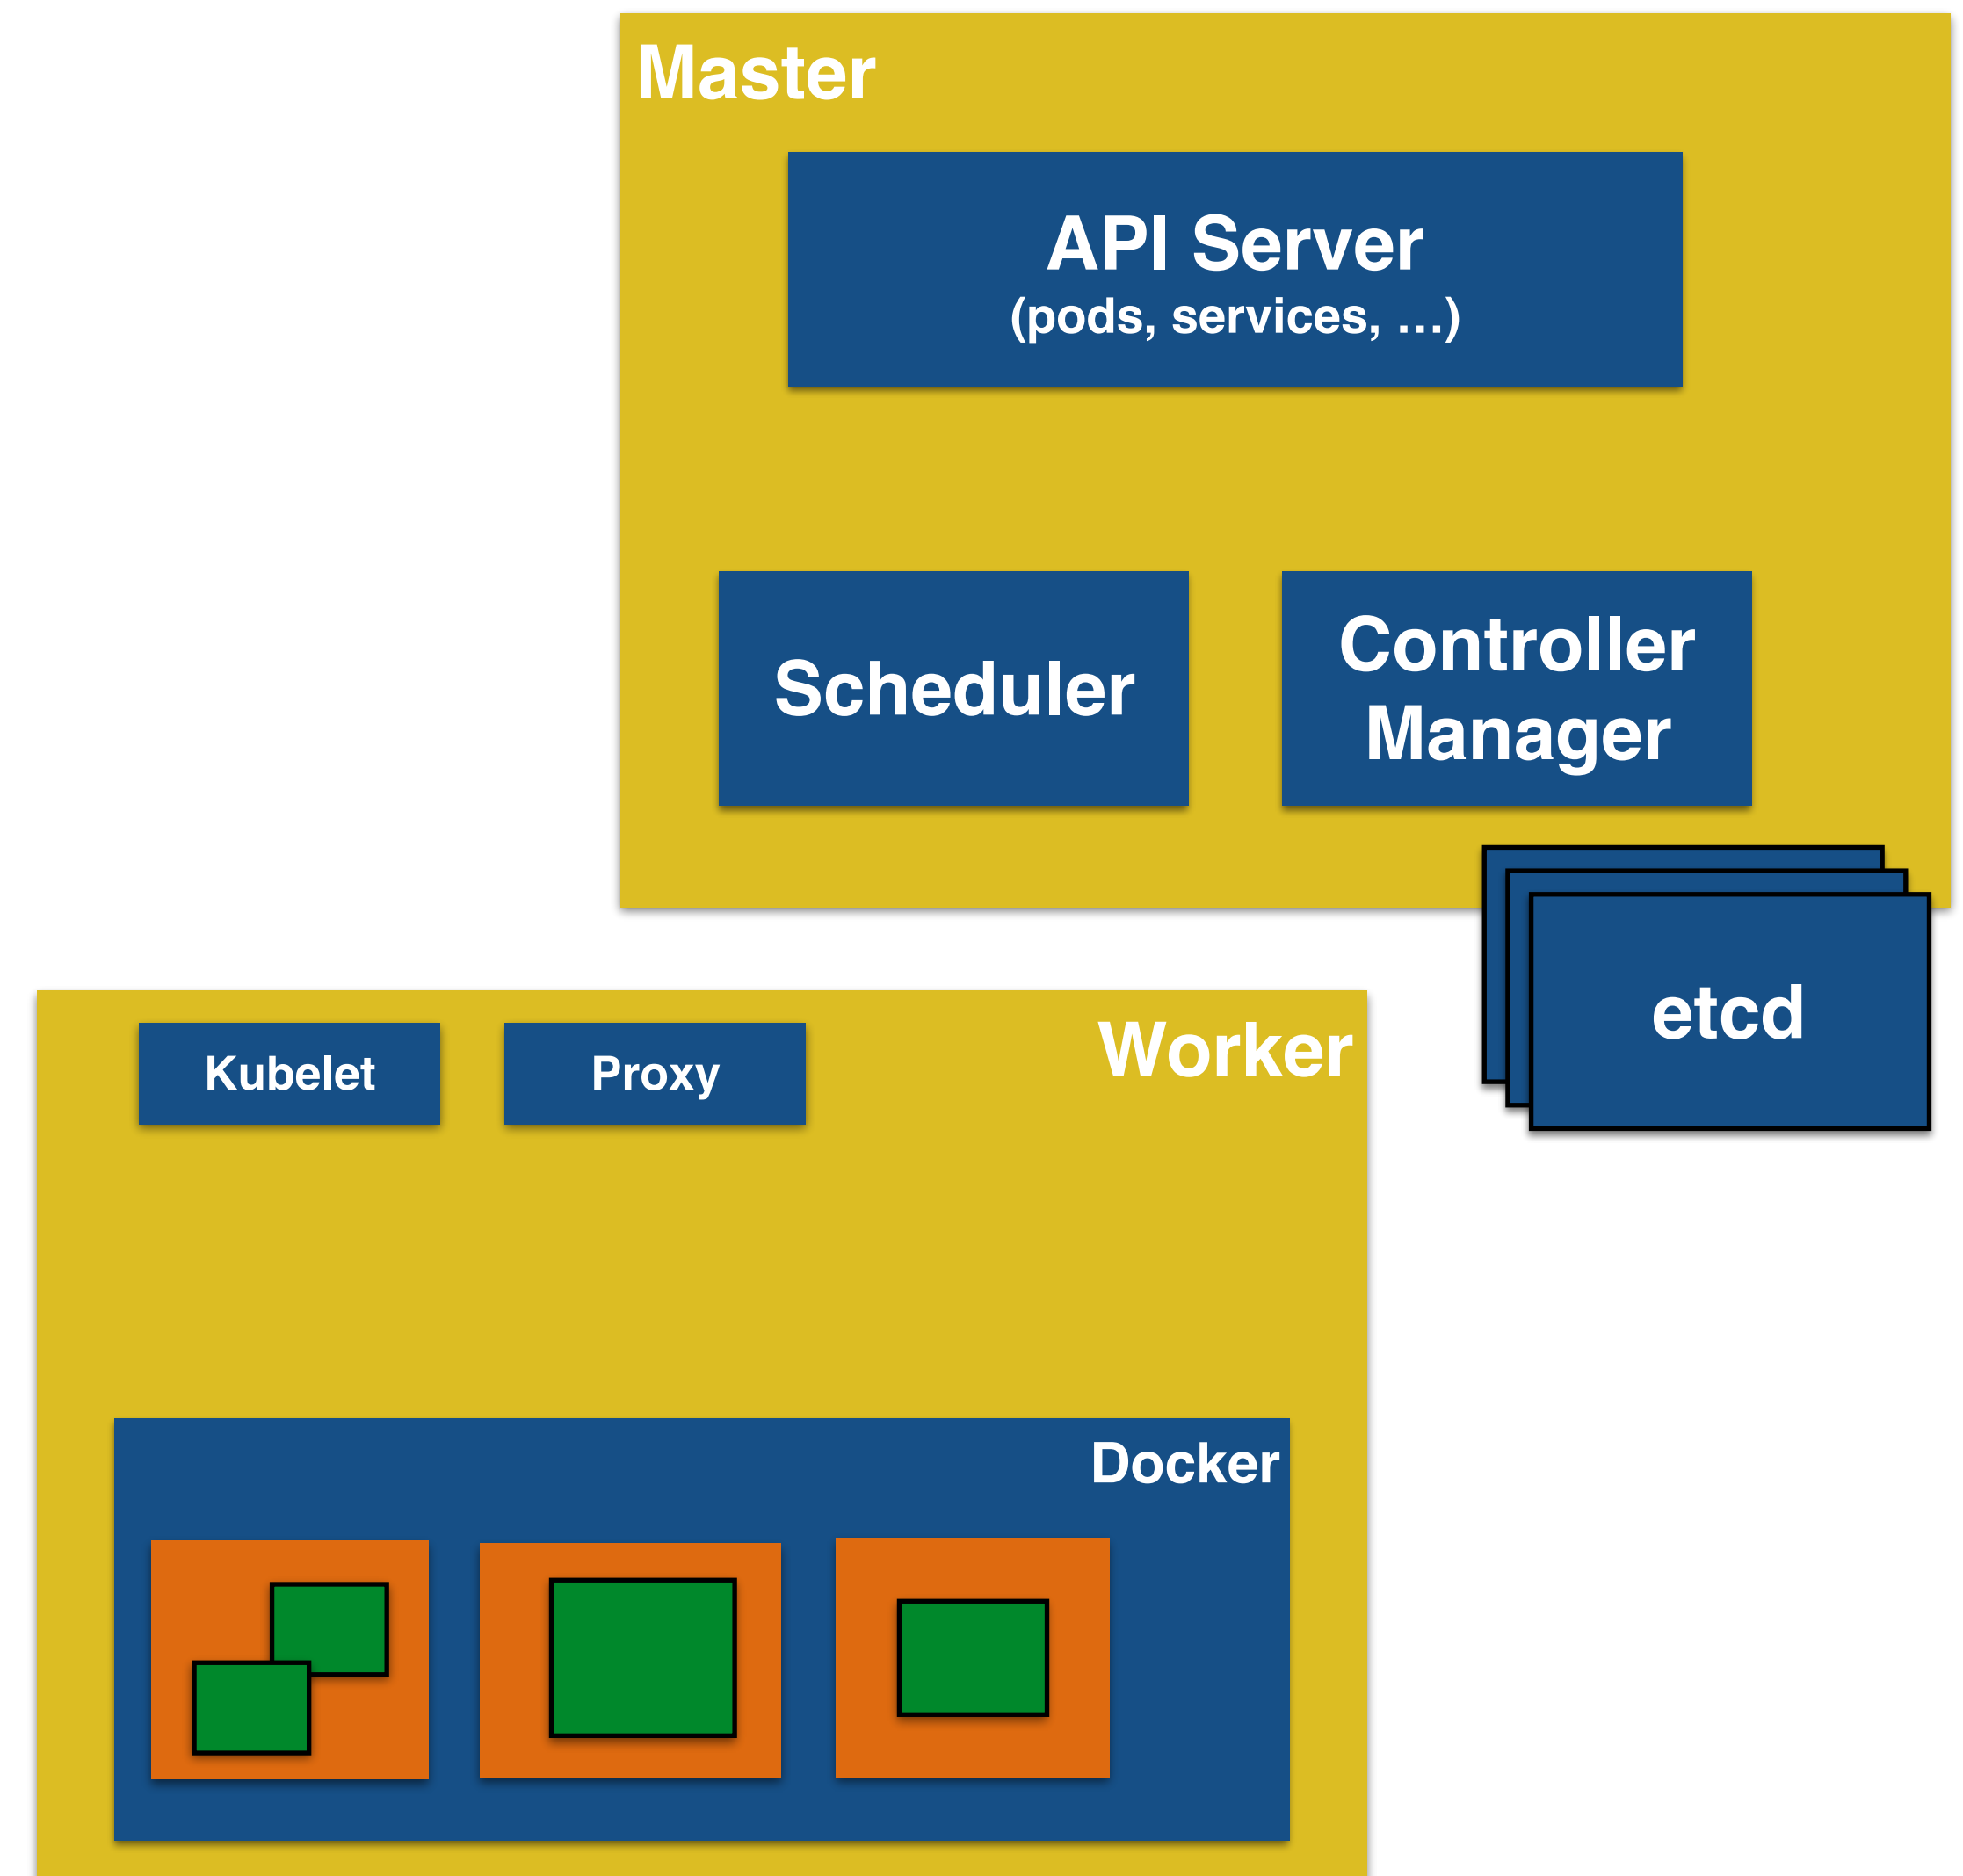
# Kubernetes Concepts

- **Pods**: collocated group of Docker containers that share an IP and storage volume
- **Service**: Single, stable name for a set of pods, also acts as LB
- **Label**: used to organize and select group of objects
- **Replica Set**: manages the lifecycle of pods and ensures specified number are running

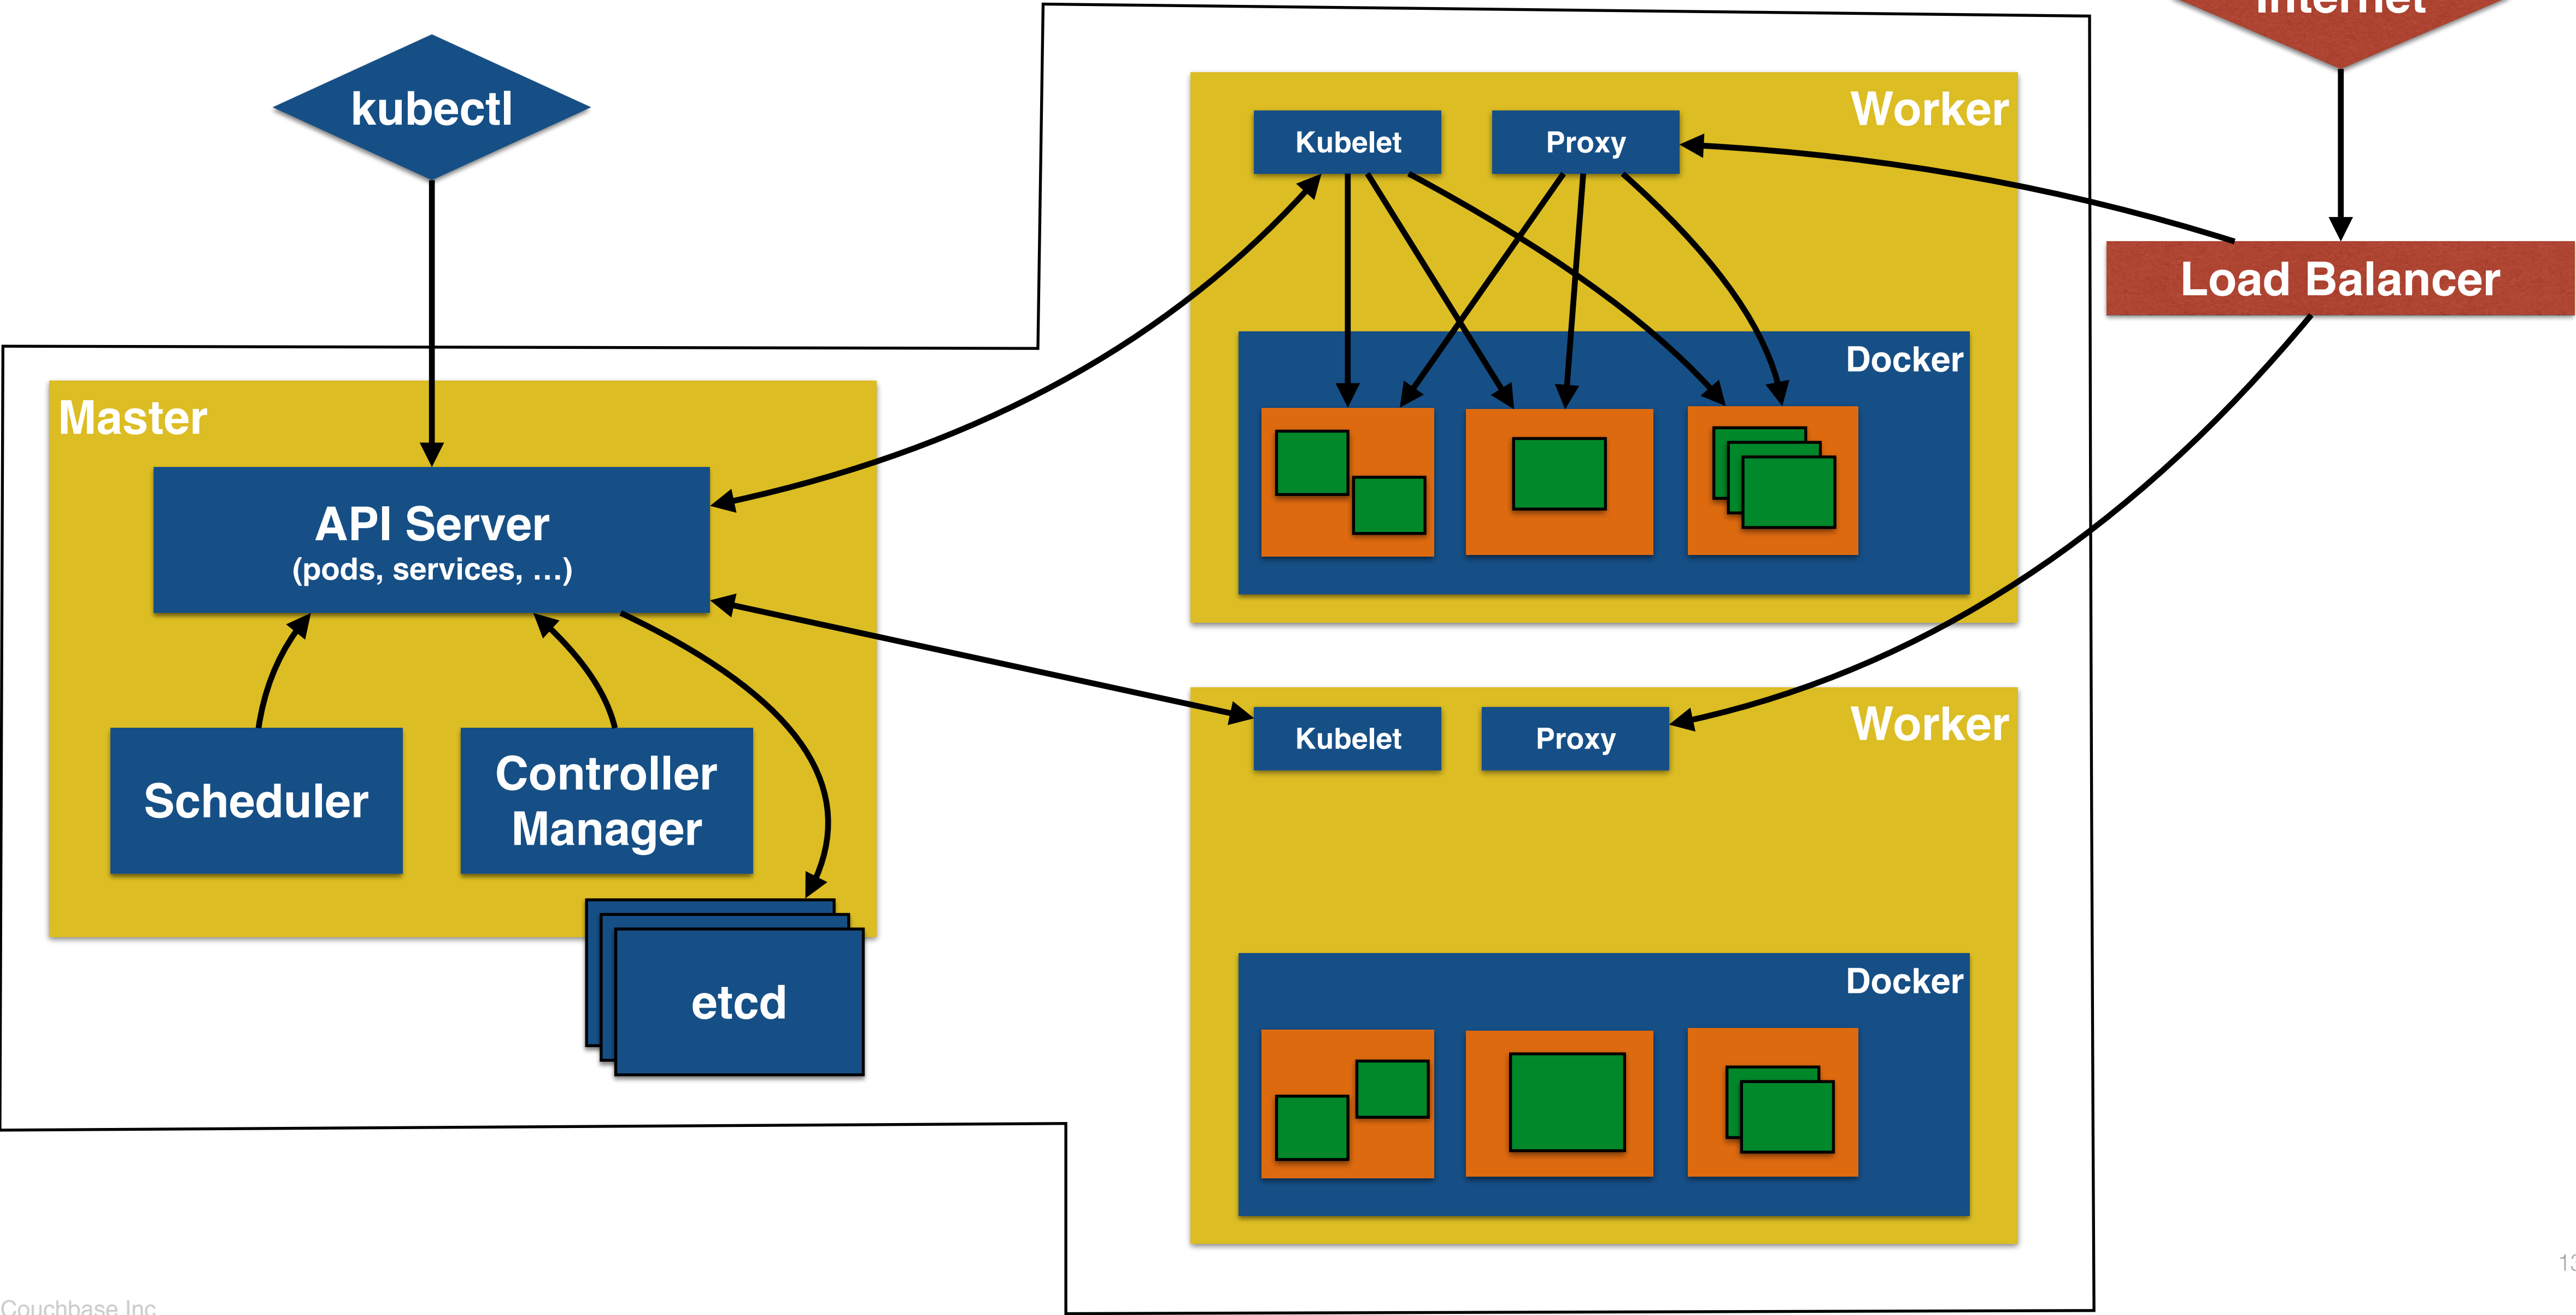


# Core Concepts: Kubernetes

- **Node:** Machine or VM in the cluster
- **Master:** Central control plane, provides unified view of the cluster
  - **etcd:** distributed key-value store used to persist Kubernetes system state
- **Worker:** Docker host running *kubelet* (node agent) and *proxy* services
  - Runs pods and containers
  - Monitored by *systemd* (CentOS) or *monit* (Debian)



# Kubernetes Cluster



# Development using Kubernetes

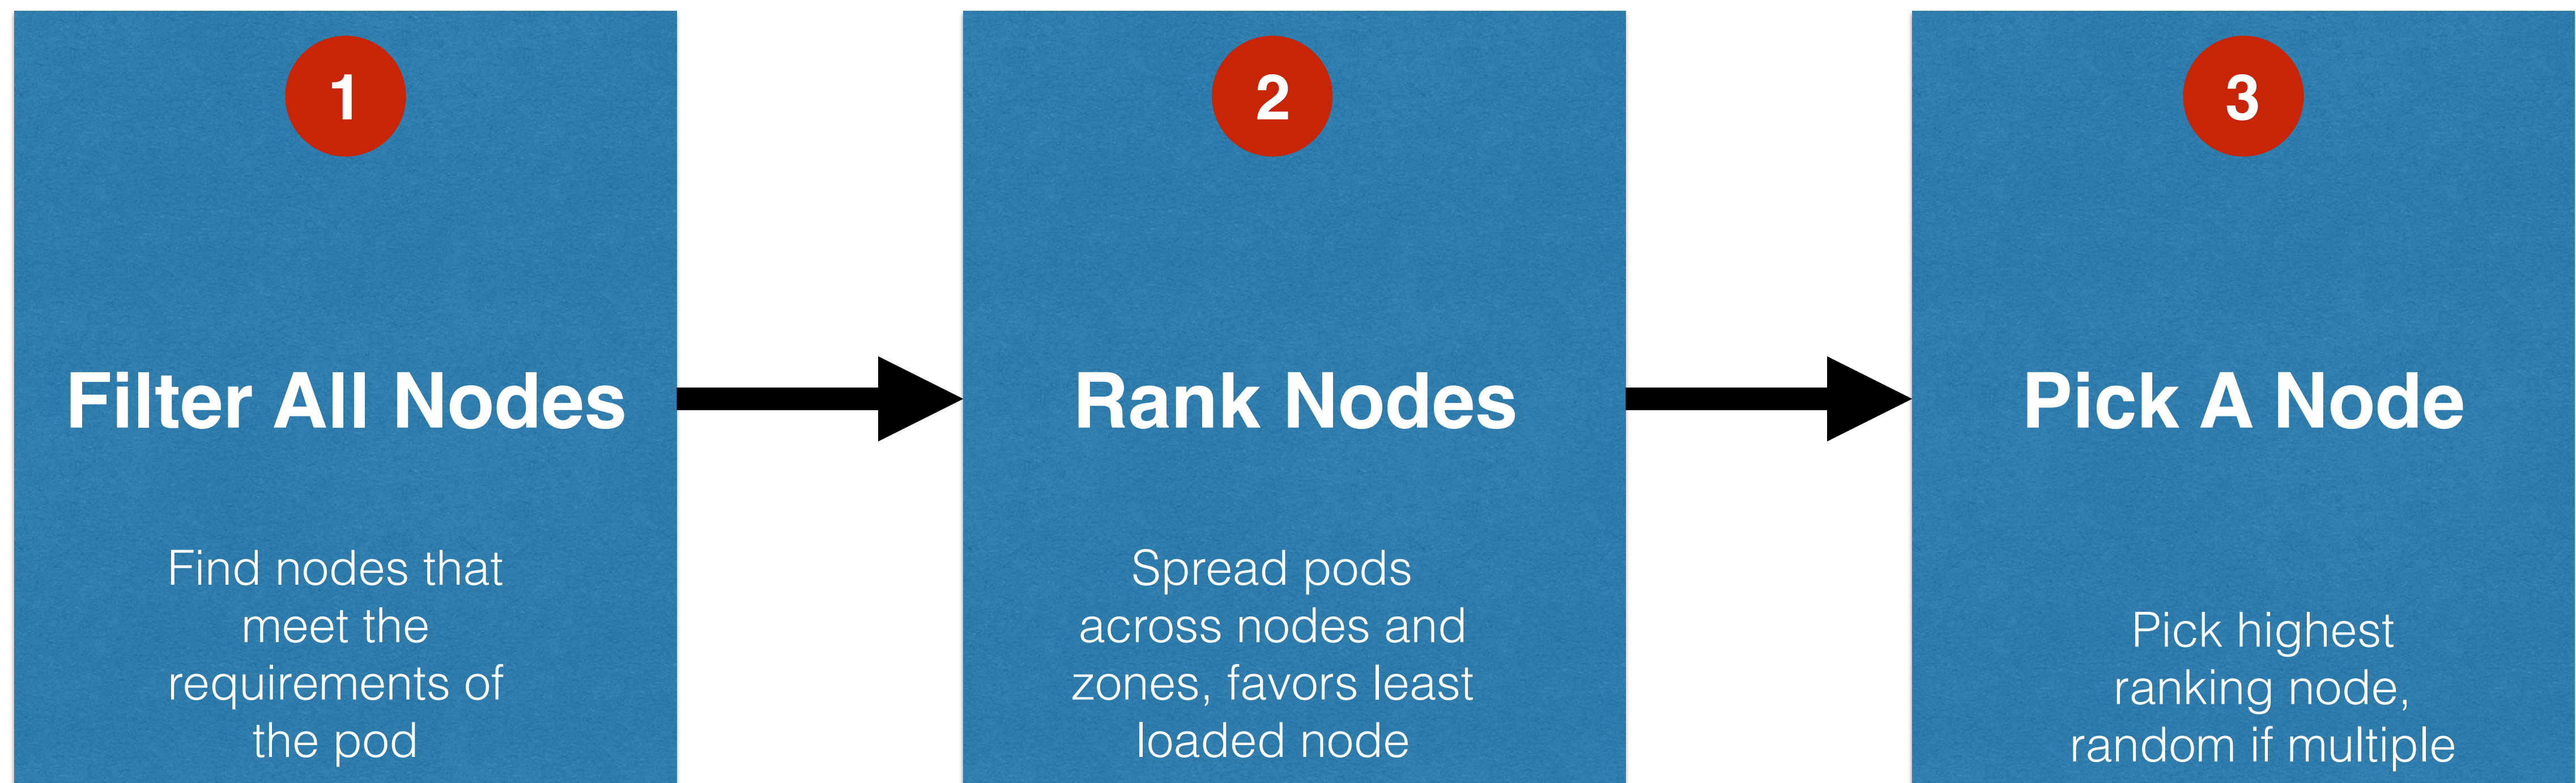
- Single node cluster
  - minikube
- Multi-node cluster
  - kops
  - kube-aws (CoreOS + AWS)
  - kube-up (deprecated)
  - Google Cloud, Azure, Tectonic, ...



# kubectl

- Controls the Kubernetes cluster manager
- `kubectl get pods or minions`
- `kubectl create -f <filename>`
- `kubectl update or delete`
- `kubectl resize --replicas=3 replicaset <name>`

# Kubernetes Scheduling Algorithm



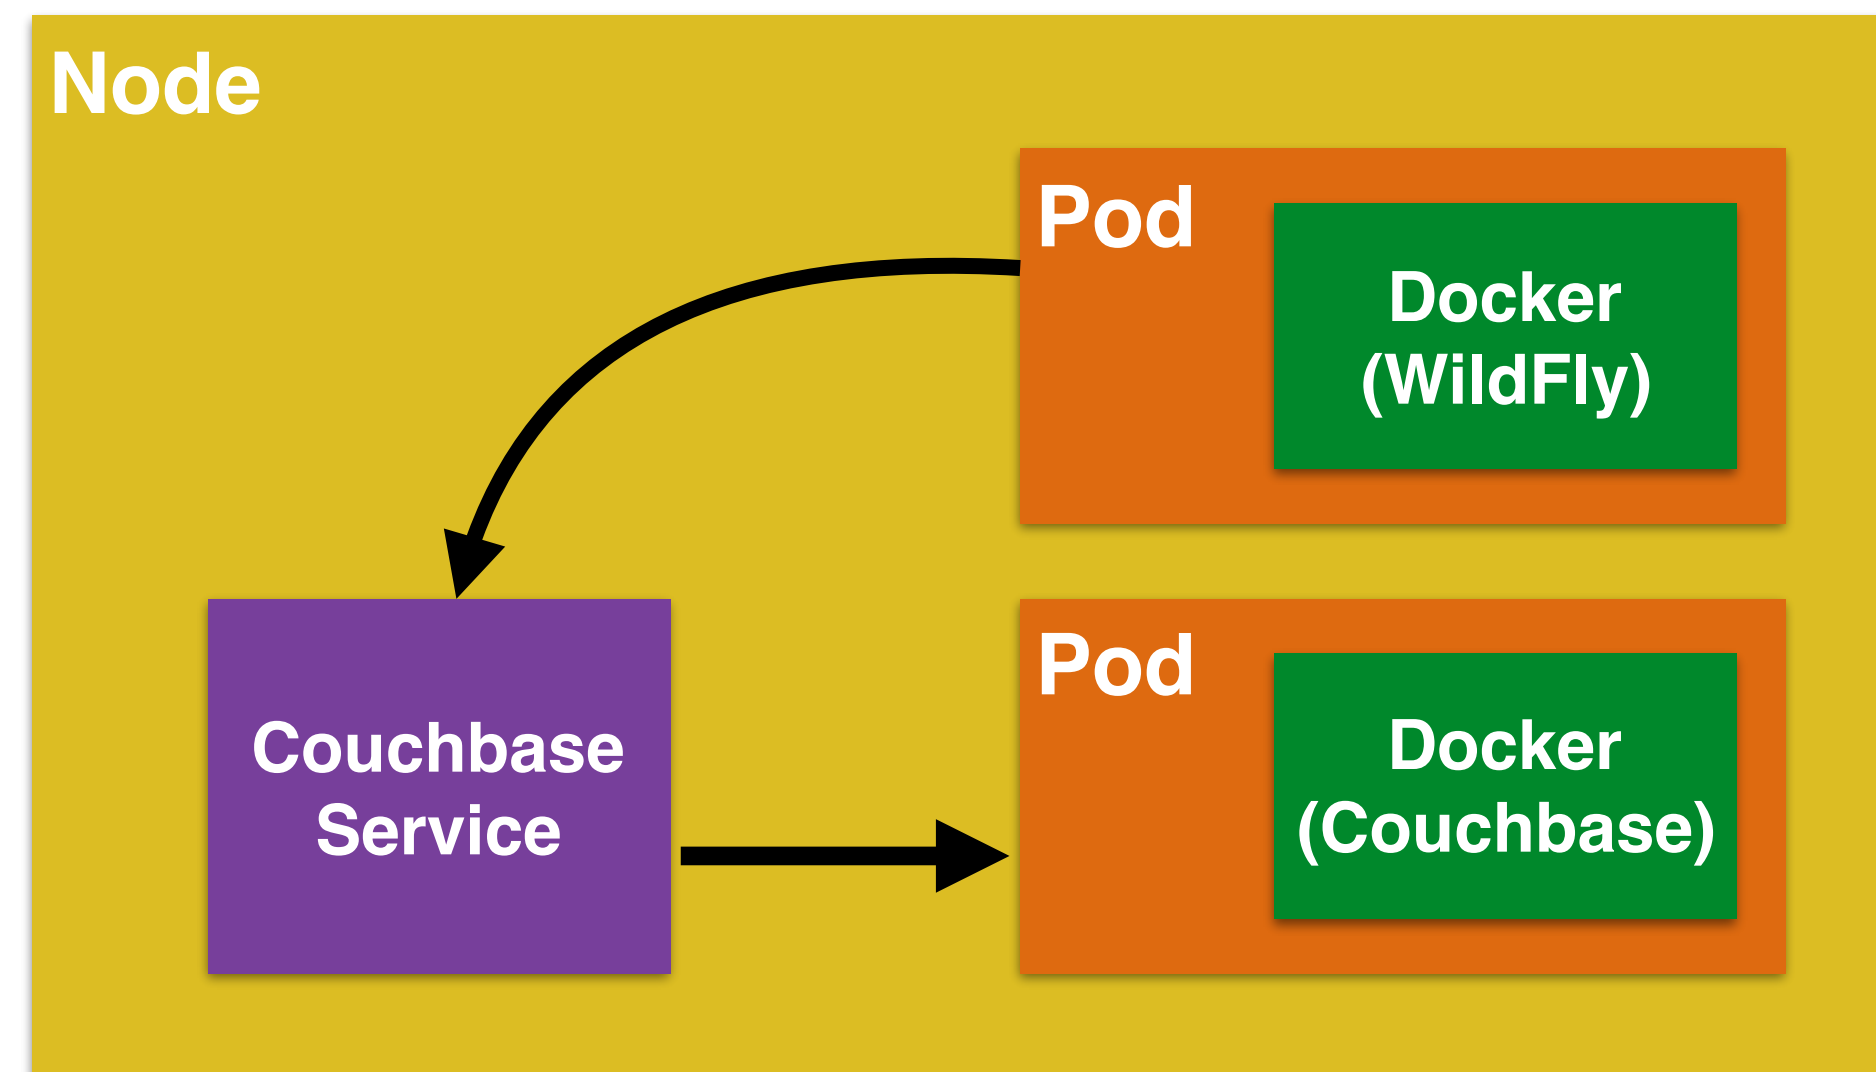


# Kubernetes Configuration File

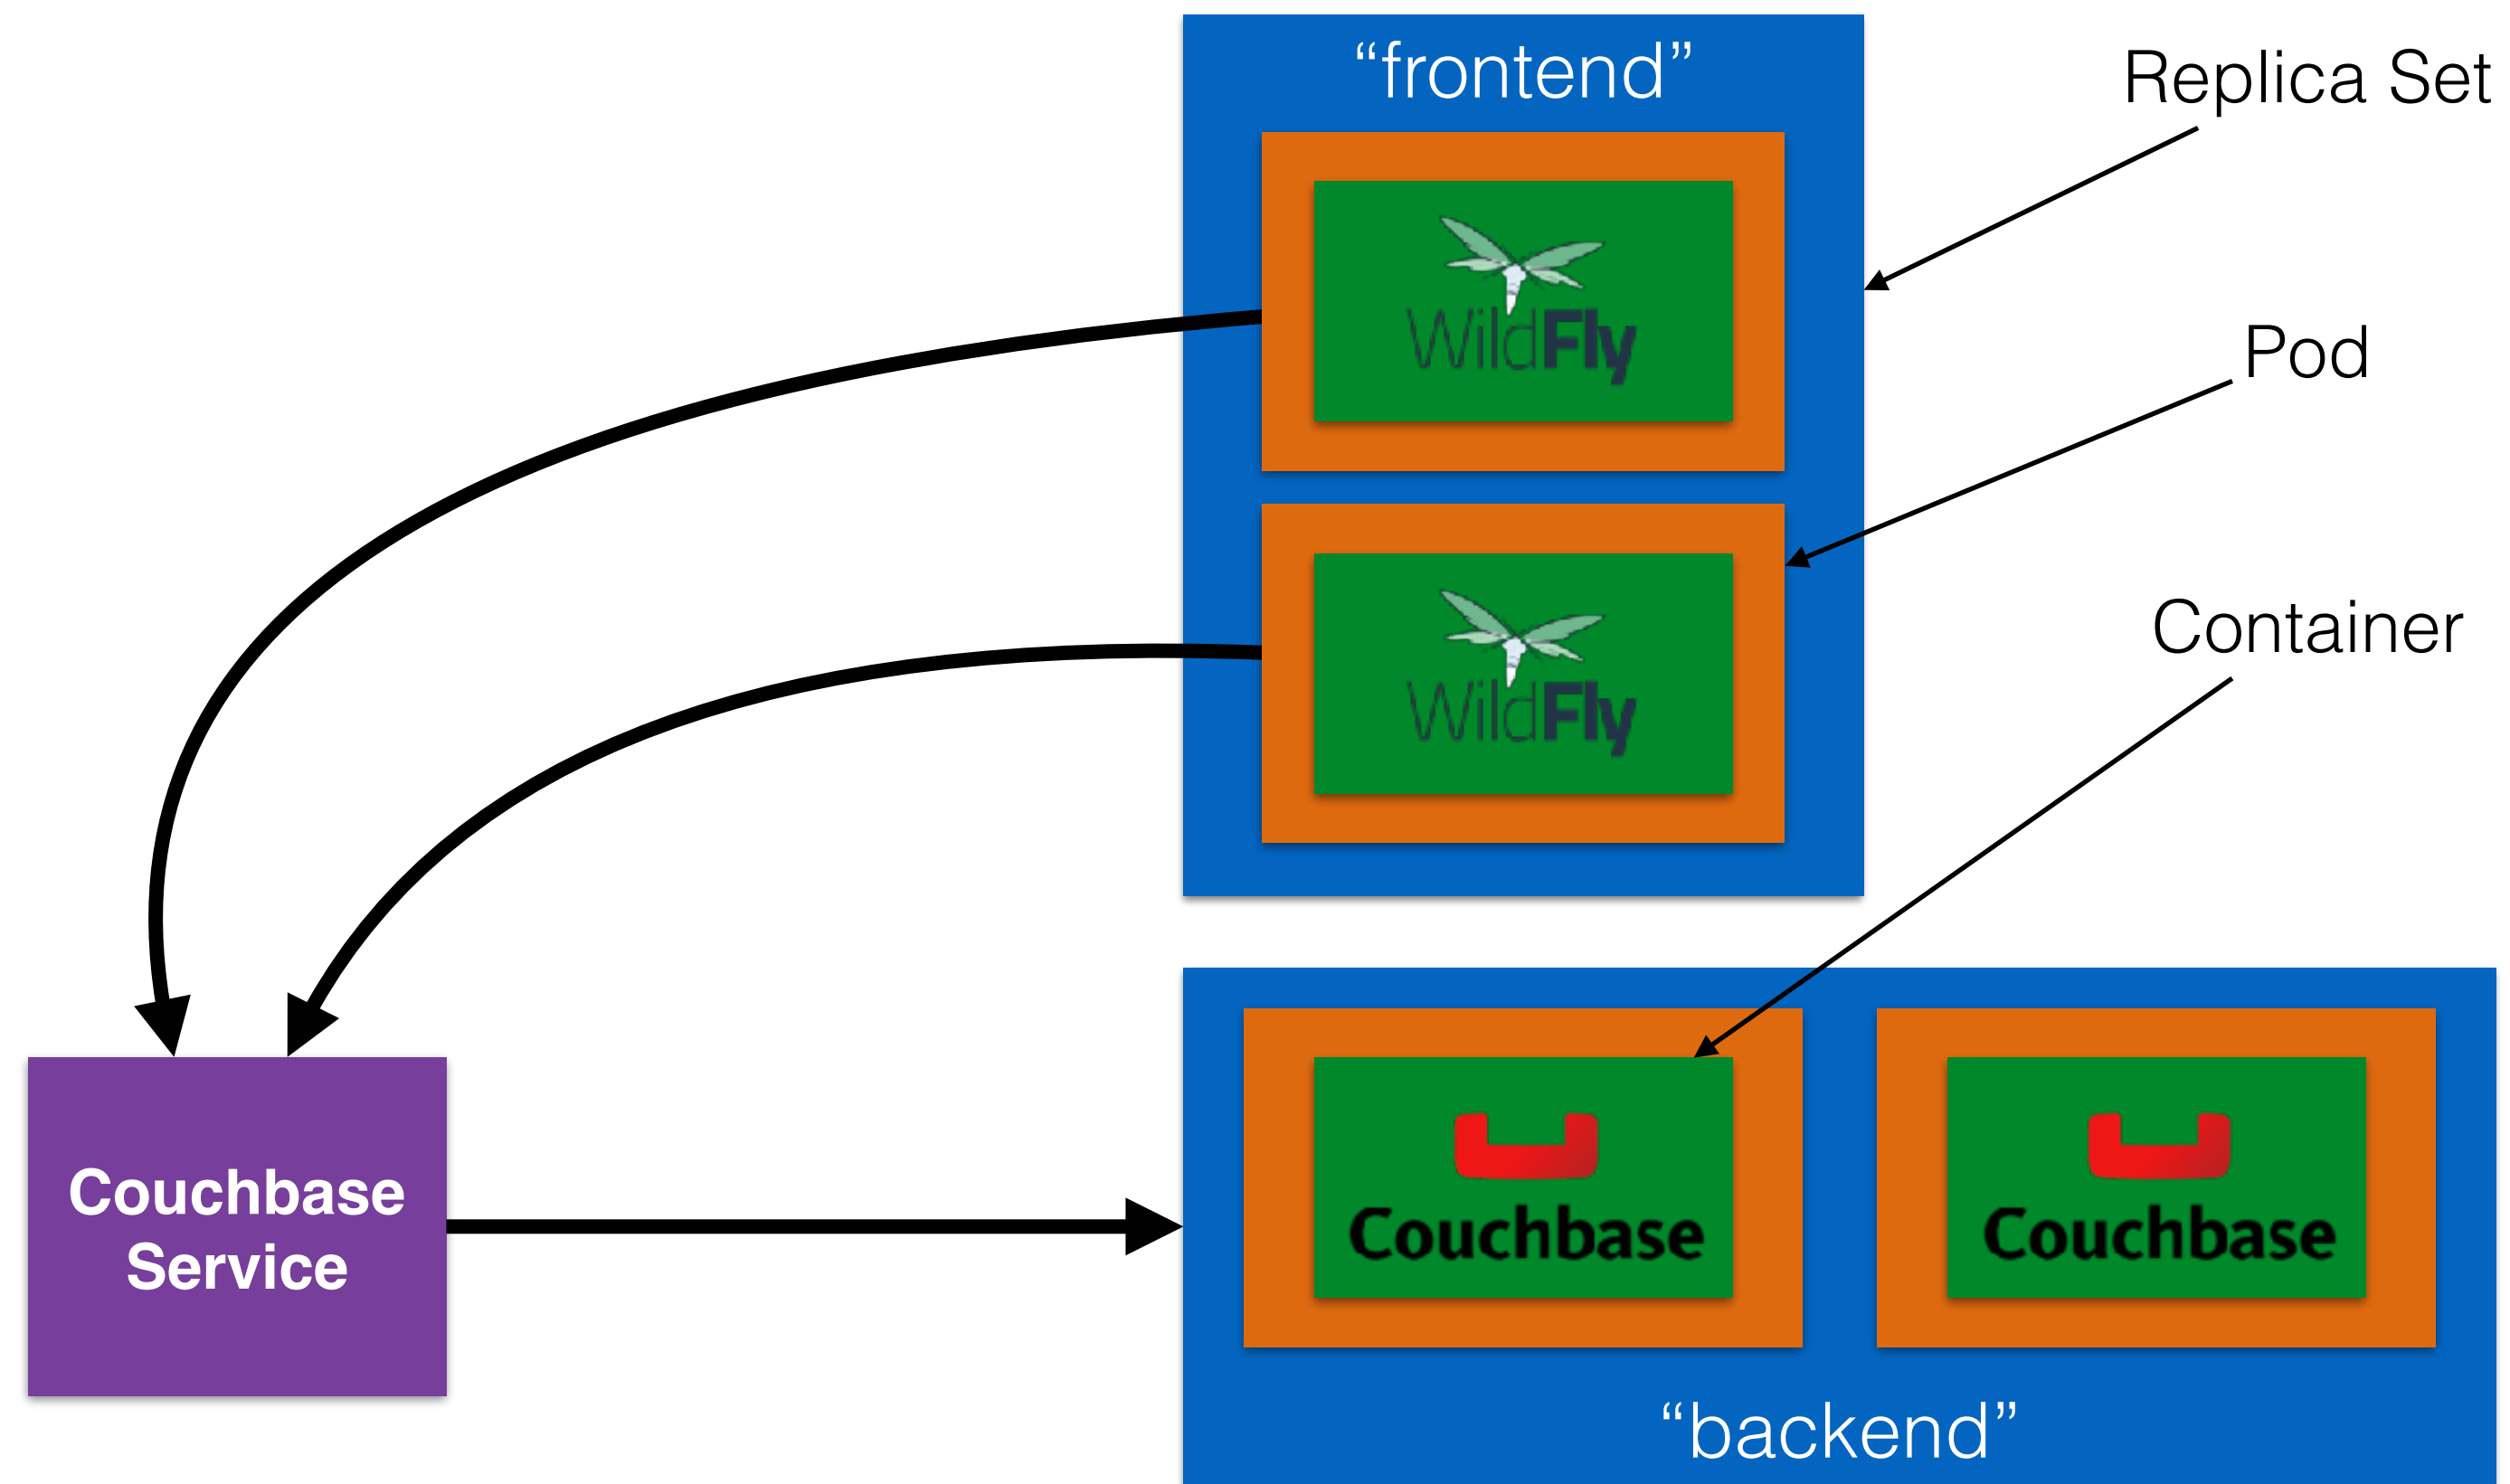
```
1  apiVersion: v1
2  kind: Service
3  metadata:
4  name: couchbase-service
5  spec:
6  selector:
7  app: couchbase-rs-pod
8  ports:
9    - name: admin
10     port: 8091
11    - name: views
12     port: 8092
13    - name: query
14     port: 8093
15    - name: memcached
16     port: 11210
17  ---
18  apiVersion: extensions/v1beta1
19  kind: ReplicaSet
20  metadata:
21  name: couchbase-rs
22  spec:
23  replicas: 1
24  template:
25    metadata:
26    labels:
27    app: couchbase-rs-pod
28    spec:
29    containers:
30    - name: couchbase
31      image: arungupta/couchbase:travel
32      ports:
33        - containerPort: 8091
34        - containerPort: 8092
35        - containerPort: 8093
36        - containerPort: 11210
37  ---
38  apiVersion: extensions/v1beta1
39  kind: ReplicaSet
40  metadata:
41  name: wildfly-rs
42  labels:
43  name: wildfly
44  spec:
45  replicas: 1
46  template:
47    metadata:
48    labels:
49    name: wildfly
50  spec:
51  containers:
52  - name: wildfly-rs-pod
53    image: arungupta/wildfly-couchbase-javaee:travel
54    env:
55    name: COUCHBASE_URI
56    value: couchbase-service
57  ports:
58  - containerPort: 8080
59  ---
```

```
graph LR
    subgraph Service [Service]
        S_name["name: couchbase-service"]
    end
    subgraph Couchbase_RS [Couchbase ReplicaSet]
        C_label["app: couchbase-rs-pod"]
    end
    subgraph Wildfly_RS [Wildfly ReplicaSet]
        W_env["name: COUCHBASE_URI<br/>value: couchbase-service"]
    end
    S_name --> C_label
    S_name --> W_env
    C_label --> W_env
```

# Services



# Services



# Mesos

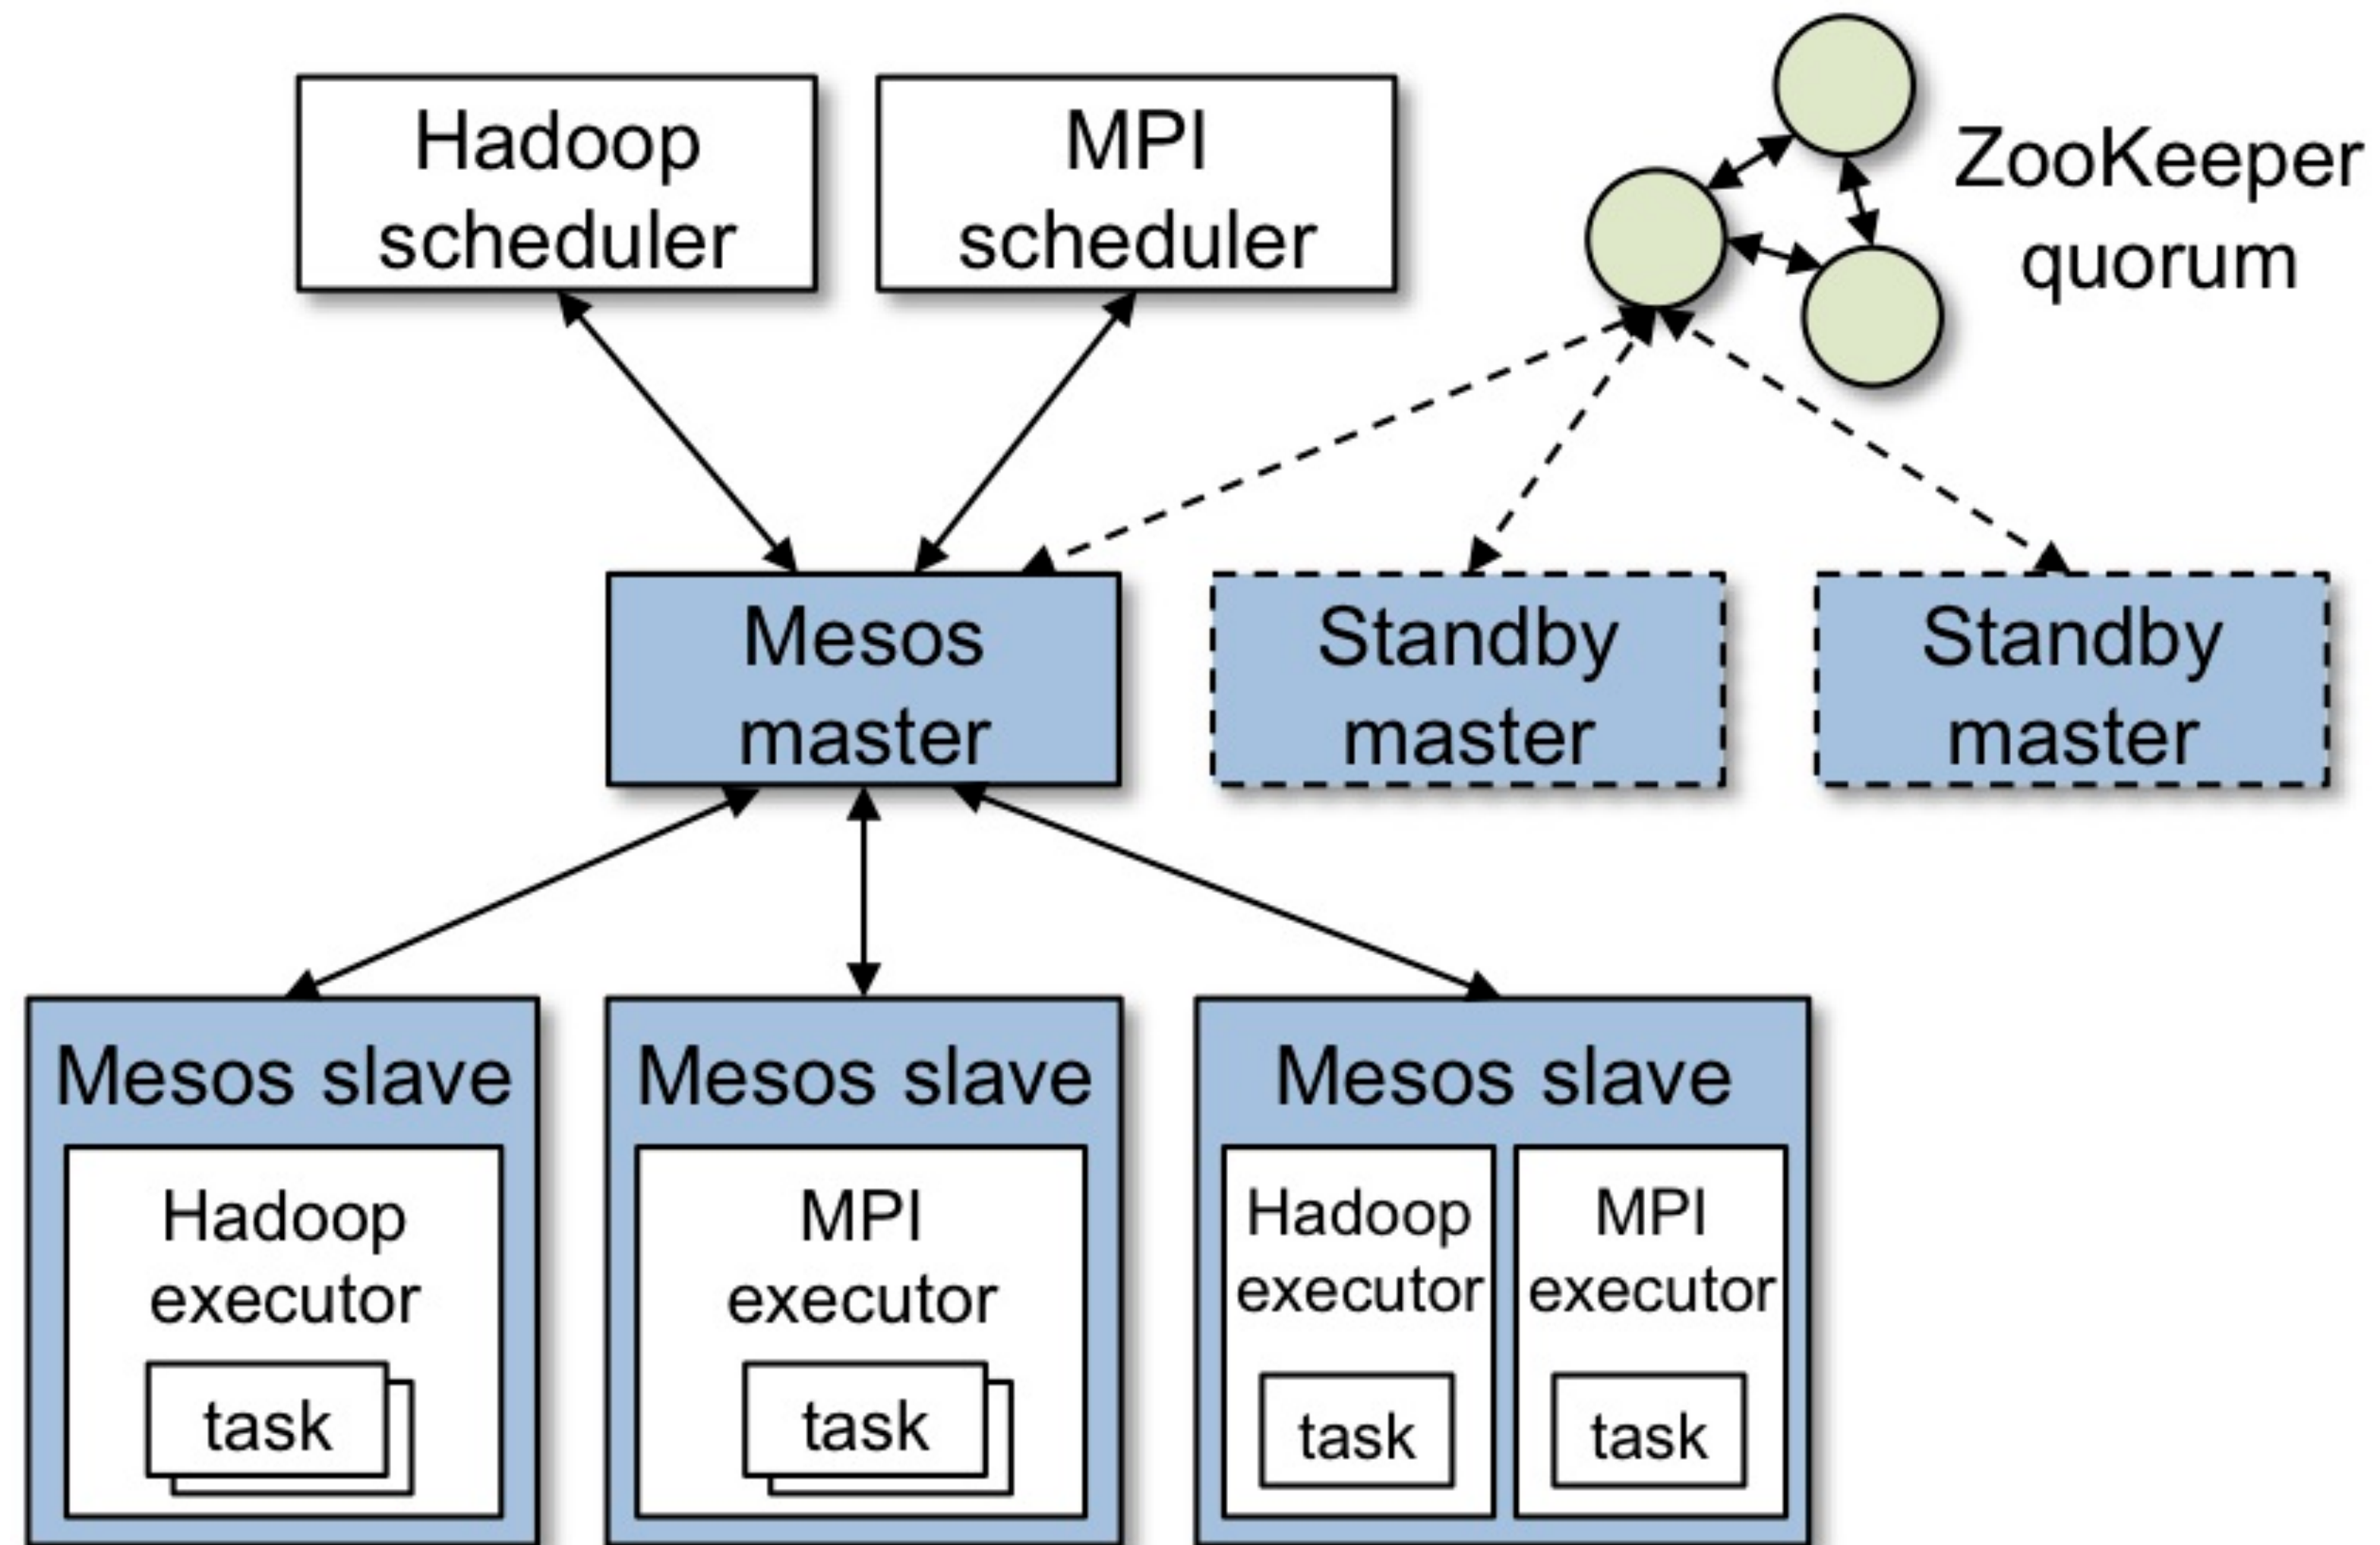
- Open source cluster manager
- Developed at UC Berkeley
- Provides resource isolation and sharing across distributed applications
- Run distributed systems on the same pool of nodes
  - Hadoop, Spark, Jenkins, Couchbase, ...
- Cluster monitoring
- Tasks isolated via Linux containers

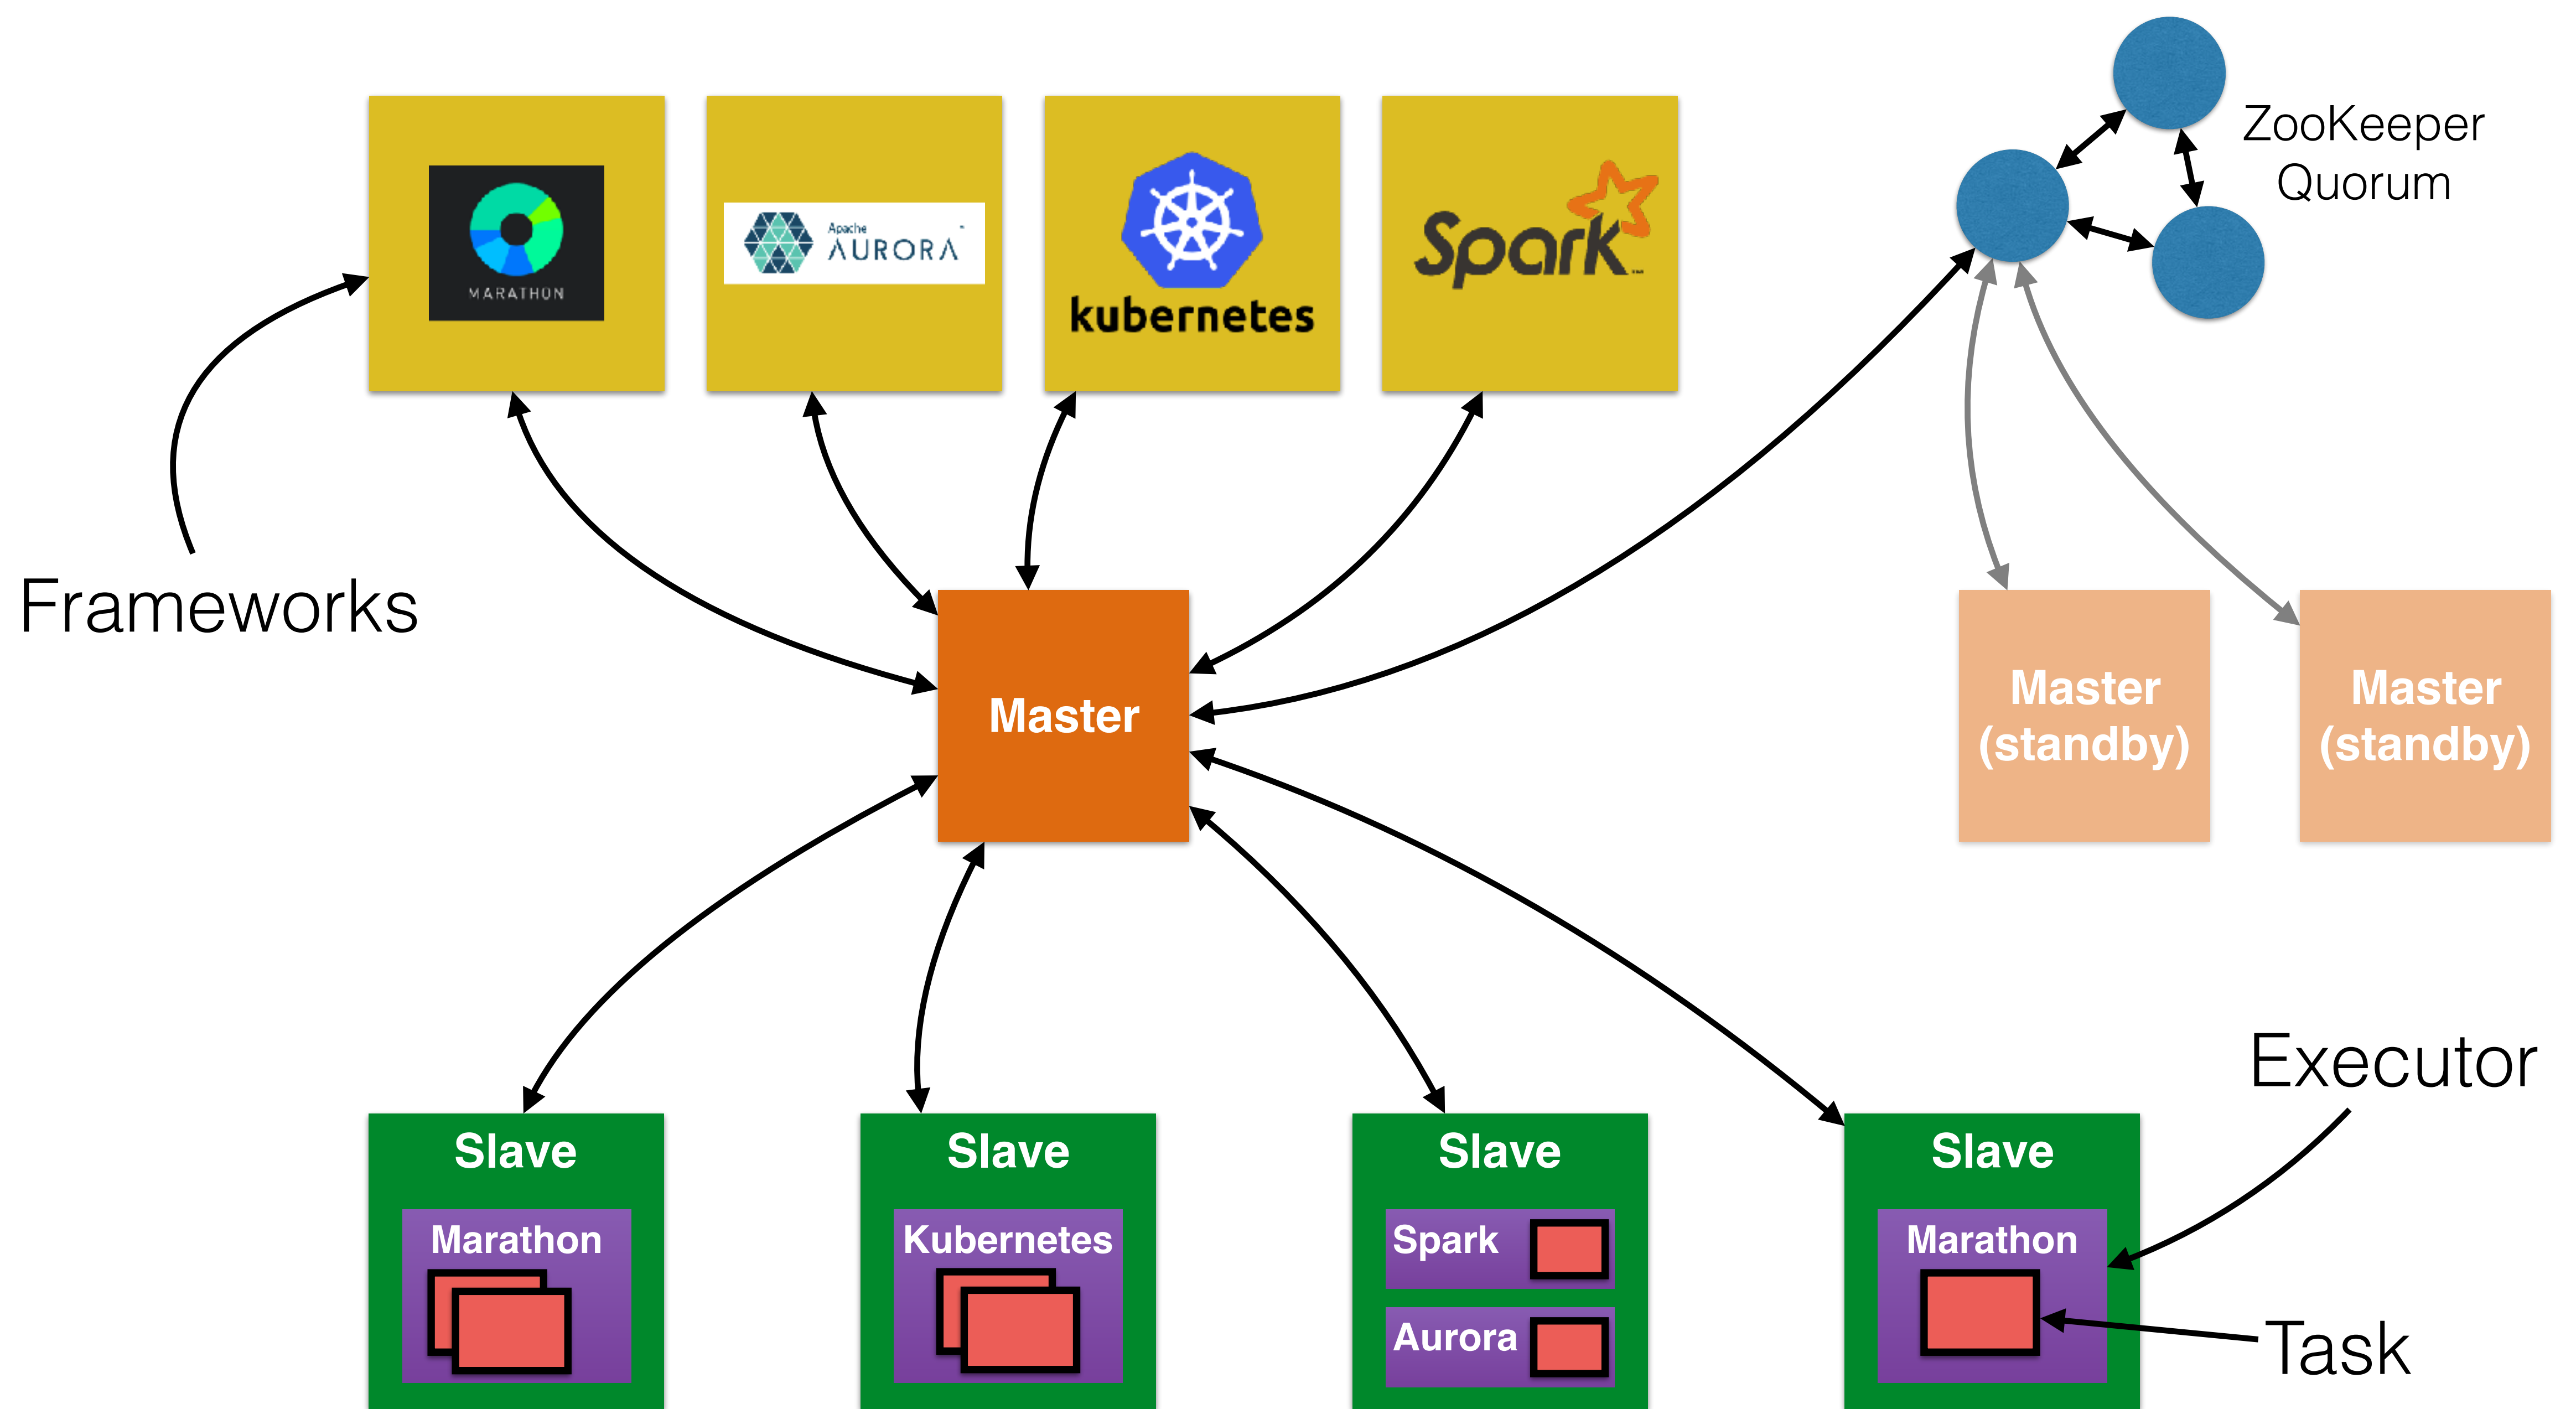


# Mesos

- Master Slave architecture
- Fault tolerant
  - Leader election via ZooKeeper
- Multi platform
  - Ubuntu, Mac OS, CentOS

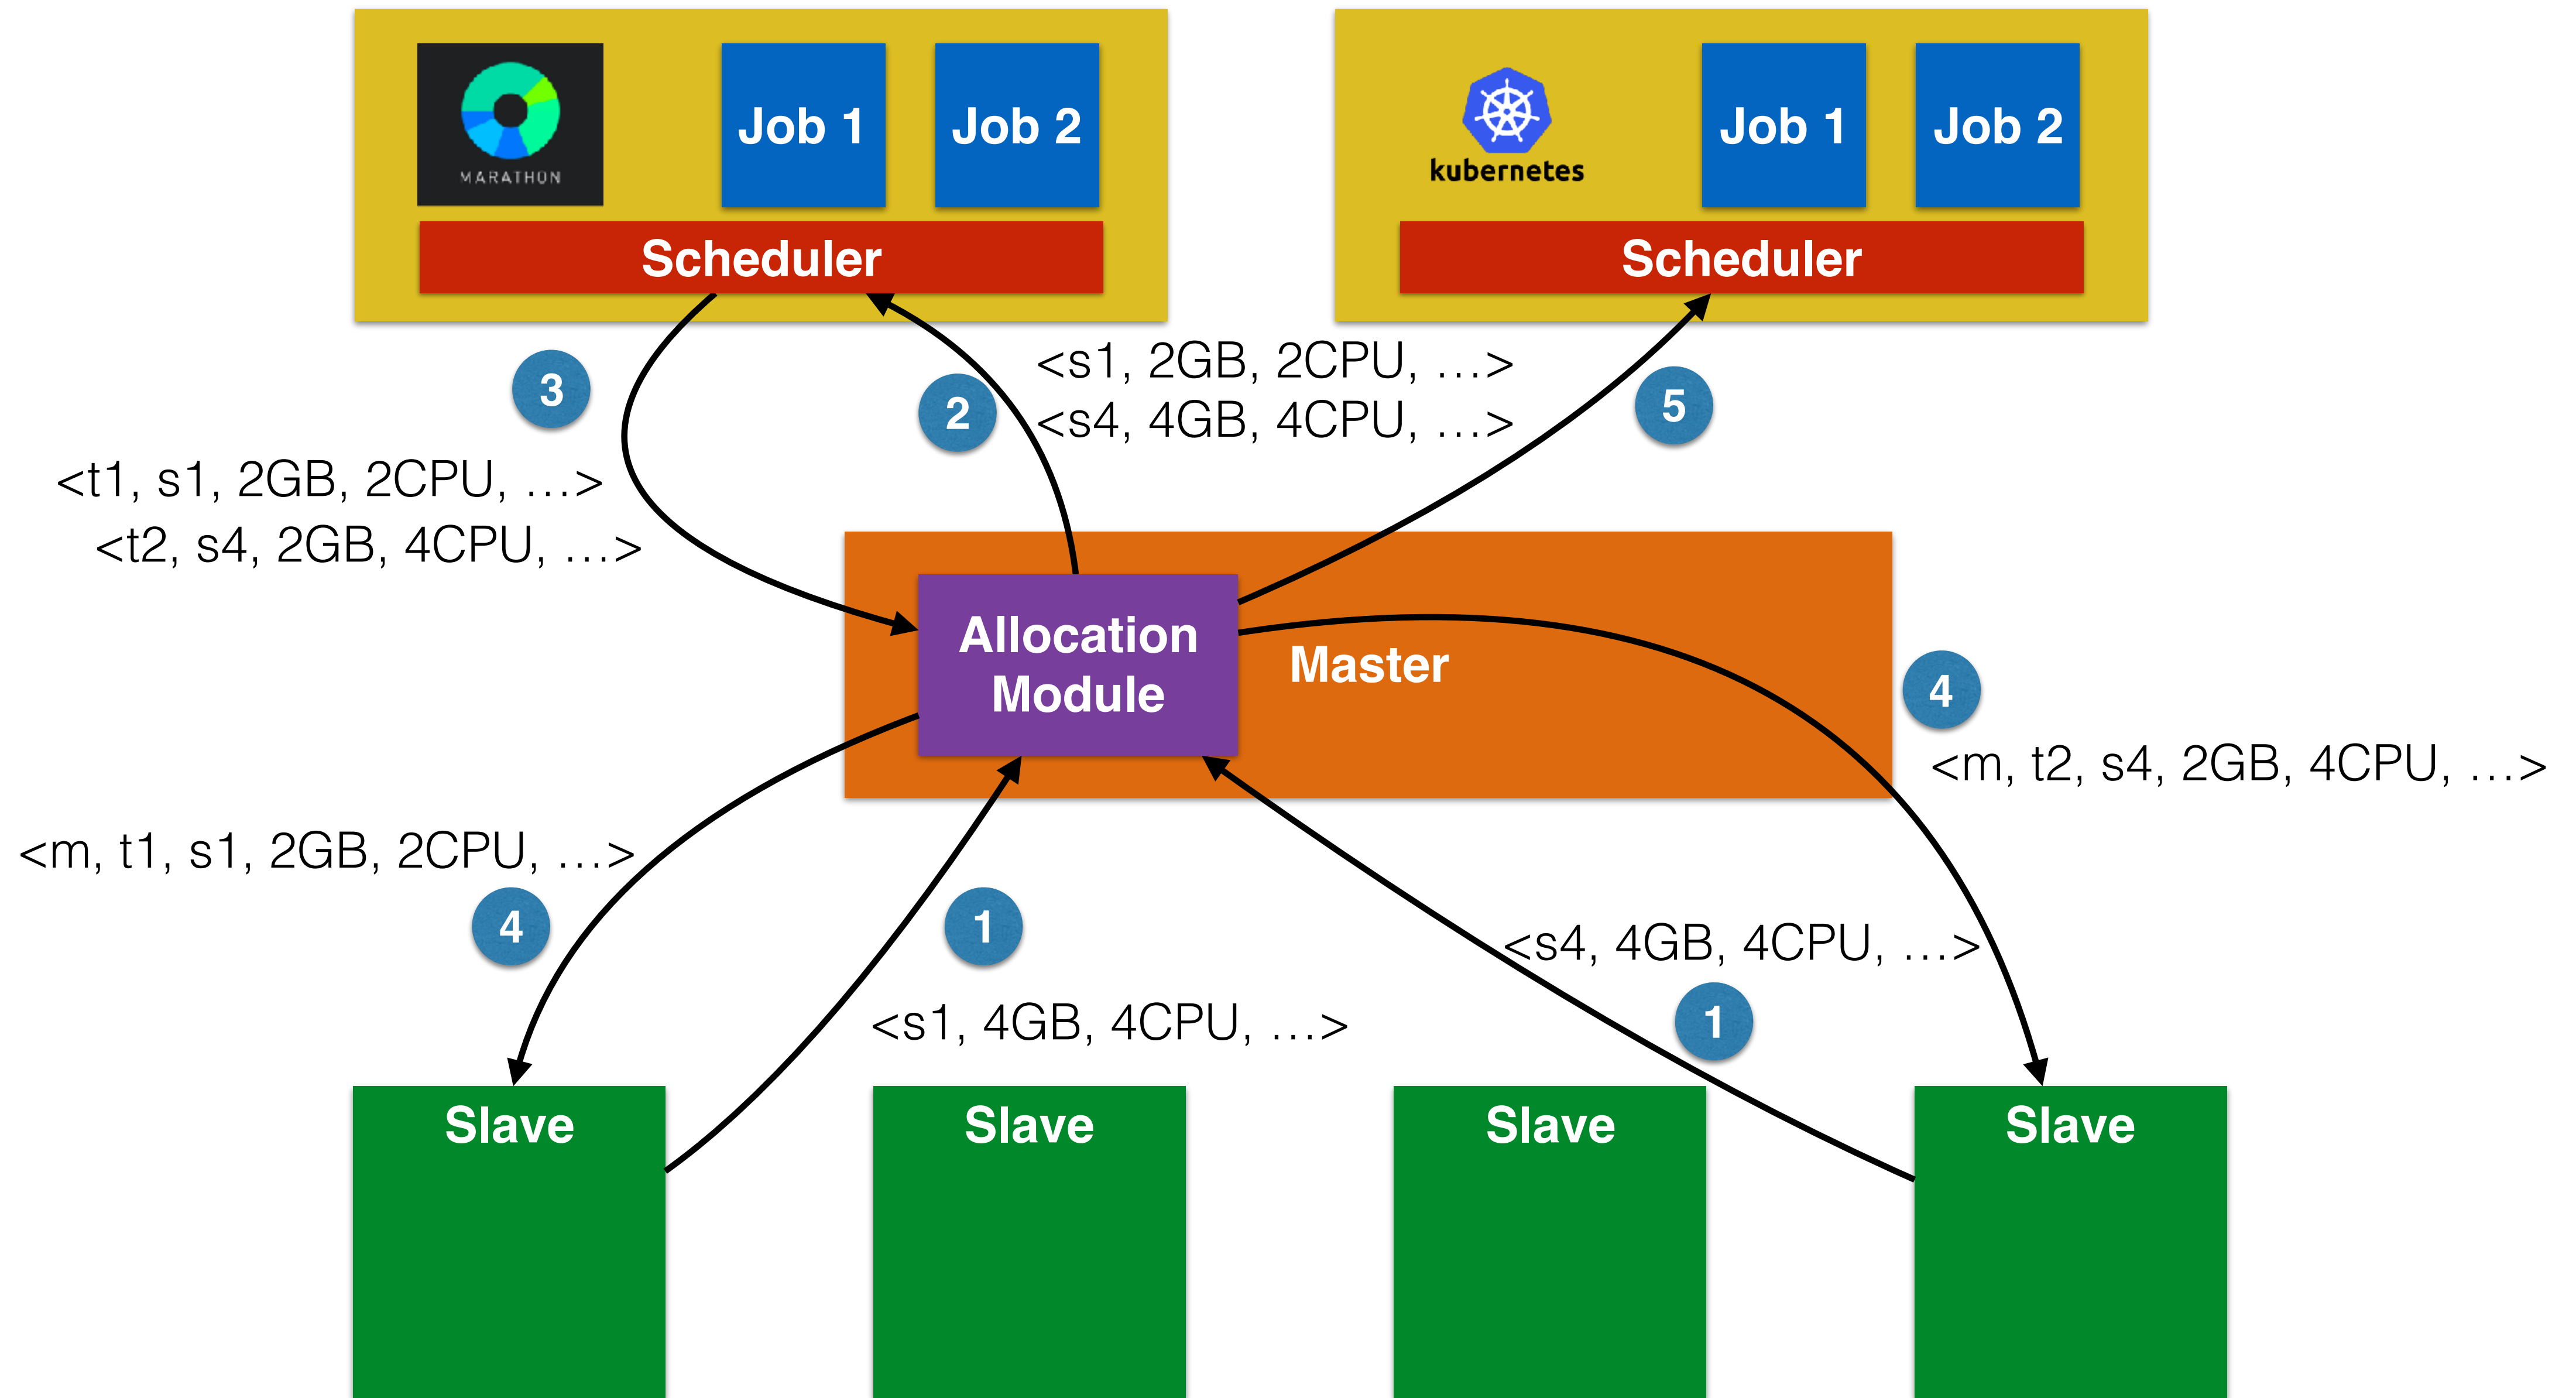
# Mesos Architecture





# Frameworks

- Frameworks are targeted at a use case and domain-specific
  - Master node “offers” resources to each framework
  - Framework “accepts” the offer and execute applications
- Framework has “scheduler” and “executor”
  - Scheduler registers with the master for “offer”
  - Executor launched on slave nodes to run the task
    - Passes a description of the task to run



# EC2 Container Service



# Thanks!

Arun Gupta, @arungupta

[github.com/javaee-samples/docker-java/tree/master/slides](https://github.com/javaee-samples/docker-java/tree/master/slides)